

EnterpriseGPT Platform

Scope & Requirements

All Features & Functionalities

Document Title:

EnterpriseGPT - Scope & Requirements

Prepared for Approval

Approval Awaited by:

Puneet Agarwal

Swati Gupta

Date: April 21, 2026

Contents

1	Core AI Reasoning & Execution Engine	2
2	Agentic AI & Orchestration	3
3	MCP & External Tool Integration	4
4	Streaming & Real-Time Communication	5
5	Knowledge & Document Intelligence	6
6	Conversational AI & Dialog Management	7
7	Voice & Multimodal Features	8
8	Frontend & User Experience	9
9	Security, Compliance & Governance	10
9.0.1	PII Protection	10
9.0.2	Access & Identity	10
9.0.3	Deployment & Data Sovereignty	10
9.0.4	Compliance & Monitoring	10
10	Analytics & Monitoring	11
10.0.1	Conversation Intelligence	11
10.0.2	Document & Knowledge Analytics	11

1 Core AI Reasoning & Execution Engine

- **Chain-of-Thought (CoT) Reasoning** — Breaks complex queries into structured, step-by-step plans with visible reasoning traces before acting
- **THINK-ACT-OBSERVE Execution Loop** — Every task passes through a planning phase, tool/action execution phase, and result evaluation phase
- **Multi-Step Planning** — Decomposes large goals into sequential subtasks that are planned and executed in order
- **Autonomous Task Execution** — Completes multi-step tasks without requiring constant user input or confirmation
- **Parallel Agent Processing** — Runs multiple subtasks simultaneously to reduce overall response and completion time
- **Proactive Assistance** — Anticipates what the user needs next and surfaces suggestions without waiting to be prompted

2 Agentic AI & Orchestration

- **Multi-Agent Orchestration** — Multiple specialized agents collaborate, share context, and hand off tasks seamlessly within a workflow
- **Central Orchestrator Model** — A master agent distributes work to specialized sub-agents within defined authority boundaries
- **Human-in-the-Loop Controls** — Configurable approval triggers require human sign-off before sensitive or high-risk actions are taken
- **Tool Calling** — Agents invoke backend APIs, databases, and real-world services as part of conversation or workflow nodes
- **Checkpoint-Based State Management** — Workflow progress and state is persisted at every step to improve multi-step continuity and recoverability
- **WorkFlow™** — A single prompt automatically sets up complete end-to-end workflow flows, and RBAC-governed Sub-Agents are auto-created and assigned to appropriate users and roles within the workflow

3 MCP & External Tool Integration

Powered by Model Context Protocol (MCP) via Composio — currently supporting 7 enterprise connectors, with 1,000+ marketplace connectors coming soon.

- **MCP Tool Integration (Composio)** — Calls external tools and services through Model Context Protocol, enabling agents to act across enterprise systems in real time
- **Gmail (Google Workspace)** — Reads, drafts, sends, and manages emails directly from within agent workflows and automation chains
- **Google Drive (Google Workspace)** — Accesses, uploads, searches, and retrieves files from Drive as part of any task or workflow execution
- **Google Calendar (Google Workspace)** — Creates, reads, updates, and manages calendar events and meeting schedules through mapped tool calls
- **Google Sheets (Google Workspace)** — Reads and writes structured data to and from Sheets within automated reporting and data workflows
- **Google Meet (Google Workspace)** — Schedules, manages, and tracks Meet sessions through automated tool calls inside agent flows
- **ServiceNow** — Creates, updates, assigns, and resolves IT service tickets and ITSM workflows directly from within agent conversations
- **Jira** — Creates issues, updates sprint boards, assigns tasks, and tracks project progress through agent-triggered Jira actions

Additional MCP connectors actively being onboarded include Slack, Salesforce, SAP, HubSpot, Zendesk, GitHub, Notion, and 1,000+ others from the Composio MCP marketplace.

4 Streaming & Real-Time Communication

- **SSE Real-Time Streaming** — Responses, plans, and intermediate step events stream live to the UI instead of waiting for a full final reply
- **Real-Time Analytics Dashboard** — Live conversation performance monitoring and usage reporting streamed to dashboards

5 Knowledge & Document Intelligence

- **Agentic RAG for Documents** — Retrieves relevant document chunks and uses them to generate grounded, cited answers
- **Document Upload & Ingestion** — Users upload PDFs and Office documents that are indexed for semantic question answering
- **Citation-Backed Responses** — Every output includes references to the exact source documents it was drawn from
- **Hallucination Control** — Built-in retrieval grounding ensures all responses remain factually accurate and source-verified
- **Permission-Preserving Search** — Respects original data access permissions so users only see documents they are authorized to access

6 Conversational AI & Dialog Management

- **Multi-Turn Dialog Management** — Supports complex, long conversations with full context retention across turns
- **Conversation History Tracking** — Chat sessions are stored so users can revisit prior prompts, steps, and responses at any time
- **Clarification Prompts** — Automatically asks follow-up questions when user intent is ambiguous or incomplete
- **Fallback Logic** — Triggers graceful fallback responses or human handoff when confidence scores are low
- **DialogGPT-Style Orchestration** — LLM-powered engine that orchestrates multi-step, human-like dialogue flows across complex conversations
- **Automatic Dialog Generation** — Builds production-ready dialog flows from just a brief intent description without manual mapping
- **Use Case Suggestions** — Auto-generates relevant enterprise use cases during the agent creation journey using LLMs
- **Training Utterance Suggestions** — Generates high-quality training utterances for each intent automatically

7 Voice & Multimodal Features

- **Voice Input/Output Support** — Users interact via speech using browser speech APIs for spoken prompts and voice responses
- **Speech-to-Text (STT)** — Converts spoken input into text using browser-native or cloud ASR models in real time
- **Text-to-Speech (TTS)** — Converts generated text answers into natural-sounding voice output
- **Multilingual Voice Support** — Handles voice interactions across 120+ languages including 14+ Indian languages

8 Frontend & User Experience

- **Dark / Light Theme Support** — Frontend includes theme switching for better usability across viewing preferences and environments
- **Step Timeline / Agent Loader UI (WorkFlow™ View)** — Live visual status indicator showing each agent step as in-progress, completed, or failed with phase and tool badges
- **SSE Live Streaming UI** — Chat responses appear token-by-token or step-by-step, giving users real-time visibility into agent activity
- **Responsive Design** — Works across desktop, tablet, and mobile browser environments

9 Security, Compliance & Governance

9.0.1 PII Protection

- **PII Shield — Redact & Restore Strategy (Split-Lane Architecture)** — PII is automatically detected, redacted from the prompt before it reaches the LLM, and then seamlessly restored in the final response using a dual-lane pipeline; the model never sees raw personal data at any point
- **PII Redaction at Model Layer** — Personal data is stripped before any prompt hits the underlying LLM, applicable across all channels and agents
- **On-Device Encrypted Processing** — For edge/offline deployments, all data is processed locally and never sent to any external cloud server

9.0.2 Access & Identity

- **Role-Based Access Control (RBAC)** — Controls who can access, build, modify, or manage each part of the platform including WorkFlow™ Sub-Agent assignments
- **Webhooks, OAuth, SSO-Ready** — Enterprise-grade identity federation and single sign-on support across identity providers

9.0.3 Deployment & Data Sovereignty

- **Private VPC / On-Premises Deployment** — Full self-hosted model hosting for maximum data sovereignty with zero data egress
- **Data Localization** — All data hosted within your own infrastructure; never leaves to unauthorized external servers
- **Zero Data Egress Architecture** — Ensures no user data, prompt, or document is sent outside the defined perimeter

9.0.4 Compliance & Monitoring

- **SOC 2 / GDPR / HIPAA Compliance** — Meets major regulatory requirements for enterprise and regulated industry deployments
- **Responsible AI Guardrails** — Automated safety controls to prevent harmful, biased, or non-compliant outputs across all agent responses
- **Audit Trails & Full Explainability** — Every interaction, step, tool call, change, and configuration is logged with timestamps for complete traceability
- **Versioning & Change Logs** — Full version history for agent configurations, prompts, and knowledge base updates
- **Governance Framework** — Platform-wide responsible AI governance covering agent behavior, output quality, and regulatory adherence

10 Analytics & Monitoring

10.0.1 Conversation Intelligence

- **Real-Time Dashboards** — Track resolution rates, fallback rates, user satisfaction scores, and error logs live across all agents
- **Root-Cause Detection** — Automatically identifies conversation failure areas and under-performing flows for targeted improvement
- **Component Monitoring** — Tracks data flow, processing times, and success rates for each individual AI component in the pipeline

10.0.2 Document & Knowledge Analytics

- **RAG Retrieval Quality Monitoring** — Tests and reports on semantic retrieval accuracy and chunk relevance scoring
- **Hallucination Rate Tracking** — Monitors how often responses deviate from grounded source documents
- **Knowledge Base Coverage Reports** — Identifies gaps in the indexed knowledge base vs. actual user queries